

<div>Optimization</div> <div>732A90 Computational Statistics Krzysztof Bartoszek (krzysztof.bartoszek@liu.se)</div> <div>24.1.2019 (P42) Department of Computer and Information Science Linköping University</div>	<div>Plan for today</div> <ul style="list-style-type: none"> • Introduction • Mathematical definition of problem • 1D optimization • kD optimization • R code examples 	<div>Optimization</div> <p>Nearly everything is optimization !</p> <ul style="list-style-type: none"> ■ Chemistry ■ Physics ■ Economics, Industry ■ Engineering <p>BUT EVEN</p> <ul style="list-style-type: none"> ■ Your mobile price plan ■ Course scheduling ■ Your lunch choice <p>STATISTICS</p> <ul style="list-style-type: none"> ■ Fit parameters to data ■ Propose optimal decision 	<div>Optimization: Example</div> <div>ANY BIOLOGICAL ORGANISM</div> <div>YOU</div>
<div>Optimization: Example</div> <div>Industry</div> <p>How to produce a cylindrical (WHY?) 0.5L beer can so it requires minimum material?</p> <p>Given a certain product minimize e.g. material usage, production effort while still meeting consumer requirements.</p>	<div>Optimization: Example</div> <div>Economics/Logistics</div> <ul style="list-style-type: none"> • Travelling Salesman Problem • Windmills • Flight schedule (especially “cheap” airlines) 	<div>Optimization: Example</div> <div>Statistics</div> <p>Maximize likelihood, model fitting</p> 	<div>Optimization: Example</div> <div>Maximal likelihood</div> <p>An i.i.d. sample (X_1, \dots, X_n) is drawn from a probability distribution $P(X \Theta)$, where Θ is an unknown parameter set.</p> <p>The joint probability of all the observations is</p> $P(X_1, \dots, X_n \Theta) = \prod_{i=1}^n P(X_i \Theta).$ <p>Find Θ that maximizes $P(X_1, \dots, X_n \Theta)$.</p>
<div>Mathematical formulation</div> <p>The goal is to minimize (maximize)</p> <p>Objective function: $f(\theta)$ (reproduction, chances of survival, quality of life, cost, profit, likelihood, fit to data)</p> <p>depending on</p> <p>Parameters or Unknowns θ (reproduction strategy, resource utilization, consumer choices, height & diameter, production, raw material choice, service times, route, flight routes/times, parameters)</p>	<div>Mathematical formulation</div> <div> $\min_{\theta \in \Theta} f(\theta) \quad \text{subject to} \quad \begin{cases} c_i(\theta) = 0, & i \in E \\ c_i(\theta) \geq 0, & i \in I \end{cases}$ </div> <p>QUESTION: What should we do if we are interested in maximization instead of minimization?</p> <p>QUESTION: What should we do if the constraints are $c_i(x) \leq 0, i \in I$?</p>	<div>Mathematical formulation</div> <ul style="list-style-type: none"> ■ Available environment ■ Volume: 0.5l of can ■ Production: Factories (F_1, F_2), retail outlets (R_1, R_2, R_3), cost of shipping $i \rightarrow j$: c_{ij}, production a_i per week, requirement b_j per week to optimize: x_{ij} amount shipped $i \rightarrow j$ per week $\min_{x_{ij}} \sum_{i,j} c_{ij} x_{ij} \quad \text{minimize shipping costs}$ $\sum_{j=1}^3 x_{ij} \leq a_i, i = 1, 2 \quad \text{production capacity}$ $\sum_{i=1}^2 x_{ij} \geq b_j, j = 1, 2, 3 \quad \text{demand}$ $\forall_{i,j} x_{ij} \geq 0$ <p>Question: What would happen if we drop demand constraint?</p> <ul style="list-style-type: none"> ■ ML: often no constraints 	<div>Mathematical formulation</div> <ul style="list-style-type: none"> • Split into pairs/triplets/quadruples • Think of some human anatomy part/organ: <ul style="list-style-type: none"> • What is its function? • What could it have been optimized for over the course of time? • Is it still under selection? • What constraints was and is it under? • Think of a situation where optimization is needed in your own student/professional/personal/financial situation. • State the problem in terms of <ul style="list-style-type: none"> • Objective function • Parameters • Constraints • Does it have a trivial solution? • 10 minutes
<div>Optimization approaches</div> <ul style="list-style-type: none"> • Constrained optimization <ul style="list-style-type: none"> ■ Lagrange multipliers, linear programming ■ E.g. LASSO ■ Not this lecture! • Unconstrained optimization <ul style="list-style-type: none"> ■ Steepest descent ■ Newton method ■ Quasi-Newton–Methods ■ Conjugate gradients <p>Why are there different methods?</p>	<div>Optimization approaches</div> <div>1D Optimization</div> <ul style="list-style-type: none"> • Function of a single parameter, find minimum • What <i>algorithm</i> would you suggest? • Golden-section search local minimum on $[A, B]$ interval (constraint) • Works by narrowing down the search interval with a constant reduction factor $1 - \alpha = \frac{\sqrt{5} - 1}{2} \approx 0.62$ <p>Question: Does α remind you of something?</p>	<div>Optimization approaches</div> <div>Golden section (minimization)</div> <pre> 1: $x_1 \leftarrow A, x_3 \leftarrow B,$ 2: while $x_1 - x_3 > \epsilon$ do 3: $a \leftarrow \alpha(x_3 - x_1)$ 4: $x_2 = x_1 + a, x_4 = x_3 - a$ 5: if $f(x_1) > f(x_2)$ then 6: $x_1 \leftarrow x_1, x_3 \leftarrow x_4$ 7: else 8: $x_1 = x_2, x_3 = x_3$ 9: end if {We know the value at 3 points!} 10: end while </pre>  <p>Wikipedia, Golden-section search</p> <p>f has to be UNIMODAL</p>	<div>Optimization approaches</div> <div>1D Optimization: Example</div> <p>732A90,ComputationalStatisticsVT2019,Lecture02codeSlide16.R</p>

Multi-dimensional optimization	General strategy	How to choose the direction?	How to choose the step size?
<p>Find</p> $\min_{\vec{x} \in \mathbb{R}^n} f(\vec{x})$ <p>Using (known, or numerically evaluated)</p> <p>Gradient $\nabla f(\vec{x}) = \left(\frac{\partial f(\vec{x})}{\partial x_1}, \dots, \frac{\partial f(\vec{x})}{\partial x_n} \right)^T$</p> <p>Hessian $\nabla^2 f(\vec{x}) = \left[\frac{\partial^2 f(\vec{x})}{\partial x_i \partial x_j} \right]_{i,j=1}^n$</p>	<ul style="list-style-type: none"> Provide a (good) starting point \vec{x}_0, $\vec{x} = \vec{x}_0$ Choose a direction \vec{p} ($\ \vec{p}\ = 1$) and step size α Move to $\vec{x} := \vec{x} + \alpha \vec{p}$ Repeat step 2 until convergence 	<p>Taylor's theorem</p> $f(\vec{x} + \alpha \vec{p}) = f(\vec{x}) + \alpha \vec{p}^T \cdot \nabla f(\vec{x}) + o(\alpha^2)$ <p>\vec{p} s.t. $\vec{p}^T \cdot \nabla f(\vec{x}) < 0$ is a <i>descent</i> direction.</p> <p>Steepest descent is</p> $\vec{p} = -(\nabla f(\vec{x})) / \ \nabla f(\vec{x})\ $	<ul style="list-style-type: none"> Expensive way: find the global minimum in direction \vec{p} Trade-off way: find a decrease which is <i>sufficient</i> <p>BACKTRACKING</p> <ol style="list-style-type: none"> Choose (large) $\alpha_0 > 0$, $\rho \in (0, 1)$, $c \in (0, 1)$, $\alpha \leftarrow \alpha_0$ repeat $\alpha \leftarrow \rho \alpha$ until $f(\vec{x} + \alpha \vec{p}) \leq f(\vec{x}) + \alpha \vec{p}^T \cdot \nabla f(\vec{x})$
Newton's method	Newton's method	Quasi-Newton methods	How to compute \mathbf{B}_{k+1} ?
<ul style="list-style-type: none"> Newton–Raphson method Hessian ignored in steepest descent If f is quadratic $f(y) = \frac{1}{2} y^T \mathbf{A} y + \vec{b}^T y + c,$ <p>then minimum</p> $\vec{p}^* = -\mathbf{A}^{-1} \vec{b}.$ <ul style="list-style-type: none"> Taylor expansion of f $f(\vec{x} + \alpha \vec{p}) = f(\vec{x}) - \alpha \vec{p}^T \cdot \nabla f(\vec{x}) + \frac{\alpha^2}{2} \vec{p}^T \nabla^2 f(\vec{x}) \vec{p} + o(\alpha^3)$ <ul style="list-style-type: none"> $x := x + \alpha \vec{p}$ where $\vec{p} = -(\nabla^2 f(\vec{x}))^{-1} \nabla f(\vec{x})$	<ul style="list-style-type: none"> $(\nabla^2 f(\vec{x}))^{-1}$ is expensive to compute, there are quicker approaches, e.g. Cholesky decomposition Hessian should be positive definite for \vec{p} to be a descent direction (if not see book) Memory expensive — need to store $O(n^2)$ elements <p>BUT</p> <ul style="list-style-type: none"> Method converges quickly esp. near optimum 	<ul style="list-style-type: none"> k iteration number Compute an approximation to the Hessian, \mathbf{B}, that will allow for efficient choice of \vec{p}. SECANT CONDITION: (quasi-Newton condition) $\mathbf{B}_{k+1}(\vec{x}_{k+1} - \vec{x}_k) = \nabla f(\vec{x}_{k+1}) - \nabla f(\vec{x}_k)$ <p>BFGS Algorithm</p> <ol style="list-style-type: none"> Choose $\mathbf{B}_0 > 0$, \vec{x}_0, $k = 0$ repeat \vec{p}_k is solution of $\mathbf{B}_k \vec{p}_k = \nabla f(\vec{x}_k)$ find suitable α_k $\vec{x}_{k+1} = \vec{x}_k + \alpha_k \vec{p}_k$ calculate \mathbf{B}_{k+1} {next slide} $k = k + 1$ until convergence of \vec{x}_k at minimum 	<ul style="list-style-type: none"> We want \mathbf{B}_{k+1} and \mathbf{B}_k to be close to each other $\min_{\mathbf{B}} \ \mathbf{B} - \mathbf{B}_k\ _{s.f.}, \mathbf{B} = \mathbf{B}^T, \text{ secant condition}$ <ul style="list-style-type: none"> $\vec{y}_k = \nabla f(\vec{x}_{k+1}) - \nabla f(\vec{x}_k)$, $\vec{s}_k = \vec{x}_{k+1} - \vec{x}_k$ $\mathbf{B}_{k+1} = \mathbf{B}_k - \frac{\mathbf{B}_k \vec{y}_k \vec{y}_k^T \mathbf{B}_k}{\vec{y}_k^T \mathbf{B}_k \vec{y}_k} + \frac{\vec{s}_k \vec{s}_k^T}{\vec{y}_k^T \vec{s}_k}$ <ul style="list-style-type: none"> Closed form Sherman–Morrison formula for \mathbf{B}_{k+1}^{-1} We have to store \mathbf{B}_k^{-1}
BFGS	Conjugate Gradient method—quadratic case	Conjugate Gradient method	Nonlinear CG method
<ul style="list-style-type: none"> BGFS: Broyden–Fletcher–Goldfarb–Shanno More iterations than Newton's method (uses approximation) Each iteration quicker, no numeric inversion Good for large scale problems Choice of \mathbf{B}_0? 	<p>Minimize</p> $f(\vec{x}) = \frac{1}{2} \vec{x}^T \mathbf{A} \vec{x} - \vec{b}^T \vec{x}$ <p>for \mathbf{A} symmetric positive definite.</p> <p>Gradient:</p> $\nabla f(\vec{x}) = \mathbf{A} \vec{x} - \vec{b} =: \vec{r}(\vec{x})$ <p>Two vectors \vec{p} and \vec{q} are conjugate with respect to \mathbf{A} if</p> $\vec{p}^T \mathbf{A} \vec{q} = 0.$ <p><i>IDEA:</i> \vec{p} and \vec{q} are orthogonal w.r.t. to an inner product associated with \mathbf{A}. Use this to find a basis that will allow for easy finding of \vec{x}.</p>	<ul style="list-style-type: none"> $\vec{p}_0 = \vec{r}_0$ $\vec{p}_{k+1} = -\vec{r}_k + \beta_{k+1} \vec{p}_k$ <p>Conjugate condition has to be satisfied so</p> $\beta_{k+1} = \frac{\vec{r}_k^T \mathbf{A} \vec{p}_{k-1}}{\vec{p}_k^T \mathbf{A} \vec{p}_k}$ <p>Exercise: check this</p> <ul style="list-style-type: none"> Convergence in $\dim(\mathbf{A})$ steps (or unless cutoff for \vec{r}_k) 	<ul style="list-style-type: none"> If $f(\cdot)$ general, use $\nabla f(\cdot)$ instead of $\vec{r}(\cdot)$ <ol style="list-style-type: none"> Choose \vec{x}_0, $\vec{p}_0 = -\nabla f(\vec{x}_0)$, $k = 0$ while $\nabla f(\vec{x}_k) \neq \vec{0}$ do find suitable α_k $\vec{x}_{k+1} = \vec{x}_k + \alpha_k \vec{p}_k$ {and now update step} $\vec{\beta}_{k+1} = (\nabla^T f(\vec{x}_{k+1}) \nabla f(\vec{x}_{k+1})) / (\nabla^T f(\vec{x}_k) \nabla f(\vec{x}_k))$ {Fletcher–Reeves update, other possible} $\vec{p}_{k+1} = -\nabla f(\vec{x}_{k+1}) + \beta_{k+1} \vec{p}_k$ $k = k + 1$ end while
Nonlinear CG method	kD Optimization: Example	kD Optimization: Example	Summary
<ul style="list-style-type: none"> Local minimum convergence But this is true of all methods that cannot “jump out” of descent path Faster than steepest descent Slower than Newton and Quasi-Newton but significantly less memory 		<p>732A90_ComputationalStatisticsVT2019Lecture02codeSlide31.R</p>	<ul style="list-style-type: none"> Optimization is everywhere Numerical methods for finding minimum 1D: Golden section (unimodal), <code>optimize()</code> kD: choose step size and direction (gradient), <code>optim()</code>